

Yue Su

📍 Pittsburgh, PA

✉ yuesu4@gmail.com

☎ 217-991-2173

in yuesu

🌐 yuesu

Summary

I work on agent systems end to end: the knowledge bases they retrieve from, the platforms they run on, the benchmark that measure them, and the tracing that shows what they cost. I shipped **Stanford SIS** at Epsilla. benchmark long-horizon agent memory and published on separating retrieval failures. I co-authored **ClawTrace**. I adapt models with SFT, LoRA, and RL, serve them behind production APIs, and ship the services around them on AWS/GCP. See more at yuesu0202.github.io [🔗](#).

Education

Carnegie Mellon University

M.S. in Artificial Intelligence Engineering

Aug 2024 – Dec 2025

GPA: 3.93

University of Illinois at Urbana-Champaign

B.S. in Mathematics; B.S. in Statistics; Minor in Computer Science

Aug 2020 – May 2024

GPA: 3.88

Experience

AI Engineer Intern

Epsilla, Inc(YC S23)

Union City, NJ(Remote)

Feb 2026 – Present

- Shipped Stanford SIS end-to-end as Next.js product over three backends. Constructed biomedical knowledge graph in **Google Spanner** with **2.5M+** articles, **Neon Postgres** for all product state, and Epsilla RAG agent runtime service. Exposed the graph to cloud agents through **MCP** servers with subgraph cache and embedding-based alert subscriptions to match user queries against newly ingested claims. Built gene monograph generation assembling ranked claims with evidence strength.
- Co-authored and released ClawTrace, an open-source cost-aware tracing platform for LLM agents. Captured every LLM call, tool call, and sub-agent spans with per-step cost and token attribution, Built AI analyst diagnoses expensive or failing runs, and distills traces into reusable agent instructions, cutting median cost by **38–64%** across benchmarks while holding quality.

Software Engineer Intern

Meituan

Shenzhen, China

Jun 2024 – Aug 2024

- Built microservice for batch PDF processing handling ingestion, processing, and delivery through an API and parallelizing work. Used **PyMuPDF** to parse text, images, layouts and CV model **Detectron2** detects document structure, improving extraction accuracy **22%**. Combined the extracted text with structure into a machine-readable document representation of typed blocks with coordinates and text to improve downstream search from structure JSON and highlighted document views for review. Containerized the services with Docker and deployed on Kubernetes.

Research

Research Assistant

Mechanical and AI lab, Carnegie Mellon University

Pittsburgh, PA

Nov 2024 – Dec 2025

- **Distributed Fault-Tolerant System:** Designed distributed fault-tolerant system with active/passive replication, Redis deduplication, and heartbeat-driven failover, detecting failures within 1s, and auto-recovering within 3s without message loss.
- **AI Agent for Scientific Exploration:** Built scientific exploration agent that analyzes 1,000+ research papers to generate evidence-backed hypotheses, integrating semantic retrieval, multi-step reasoning, self-critique, human feedback, and Google Scholar API search, reducing prior-work overlap and weakly supported hypotheses by 40%.
- **Meta-Learning vs Finetune on Protein Mutation Sequence:** Developed a PyTorch MAML framework for multi-task protein mutation sequence prediction with 16-shot adaptation, modular training across multiple protein datasets, and input sequence encoding, achieving 46% lower MSE, 10% higher accuracy, and 73% lower training time than finetuning baselines.
- **Generative AI-Powered Document Automation:** Built end-to-end platform that converts PDFs into structured PowerPoint slides and narrated videos, with Python REST APIs React frontend for upload/status tracking, and AWS Lambda, EC2, S3, and SQS orchestration to reduce processing time to under 5 minutes per document and compute cost by 35%.

Publications

ClawTrace: Cost-Aware Tracing for LLM Agent Skill Distillation [🔗](#) *ACM CAIS 2026 Workshop on Agent Skills (Oral)*
Diagnosing Retrieval vs. Utilization Bottlenecks in LLM Agent Memory [🔗](#) *ICLR 2026 Workshop on MemAgents*

Skills

Languages: Python, C++, Java, SQL, JavaScript/TypeScript, Bash

Post-training: PyTorch, SFT, LoRA, RL, RLHF, Hugging Face

AI/ML: , Transformers, RAG, LangGraph, LangChain

Backend & Systems: FastAPI, Next.js, REST APIs, LLM APIs (OpenAI, Claude, Gemini), Distributed Systems, React, Spark

Cloud & Infrastructure: AWS, GCP, Docker, Kubernetes, CI/CD, Git, Linux